

IAEM Doctoral School, UL, France
LARSEN Team, LORIA, UMR 7503, France

Robotic Coverage and Exploration Tasks as Sequential Decision-Making Problems

December 12, 2017

Presented by Nassim Kaldé

Supervisors: Dr. François Charpillet (LORIA/Inria Nancy)
Pr. Olivier Simonin (CITI/INSA Lyon)

Reviewers: Pr. Philippe Mathieu (LIFL/Lille 1)
Pr. René Mandiau (LAMIH/UVHC)

Examiners: Dr. Laëtitia Matignon (LIRIS/Lyon 1)
Dr. Aurélie Beynier (LIP6/UPMC)

Contents

- 1 Motivation: coverage, exploration, and decision making
- 2 Background: robotics and graph theory
- 3 Study 1: modeling coverage
- 4 Study 2: optimal coverage
- 5 Study 3: modeling exploration
- 6 Study 4: greedy exploration
- 7 Conclusion

Coverage and exploration tasks

Coverage and exploration

- ▶ Coverage: robots visit a **known** environment to fully observe it
- ▶ Exploration: robots visit an **unknown** environment to fully map it

Recent applications (Robotics)

- ▶ NASA Mars rovers: Spirit (2003), Opportunity and Curiosity (2012)
- ▶ DARPA Robotics challenge: search and rescue, decommissioning (2012)
- ▶ iRobot Roomba vacuums: millions of robots sold (since 2002)



Decision making in Artificial Intelligence

Sequential Decision-Making (SDM)

- ▶ Formalisms: specify worlds with agents facing a series of choices
- ▶ Planners: optimize decisions of the agents to fulfill an objective

Recent milestones (Game Playing)

- ▶ IBM DeepBlue vs World Chess Champion Kasparov (1997)
- ▶ DeepMind AlphaGo vs World Go Champion Sedol (2016)
- ▶ CPRG DeepStack vs HUNL Poker pros (2017)



Bridging coverage, exploration, and decision making

Illustration with robot vacuums

- ▶ First time: build a map while cleaning (exploration)
- ▶ Then: clean efficiently based on the map (coverage)



Coverage and exploration planning

- ▶ Where to go next? Information gathering
- ▶ We model coverage and exploration toy domains!

Robotics: active sensing

Control

Problem	Question
Motion Planning [LaValle, 2006]	How to get from A to B?

State estimation

Problem	Question
Localization [Thrun et al., 2001]	Where am I?
Mapping [Elfes, 1989]	Which environment is it?
SLAM [Durrant-Whyte and Bailey, 2006]	Where am I and in which environment?

Active Sensing: control and estimation [Bajcsy, 1988, Mihaylova et al., 2003]

Problem	Where to go next to improve pose estimates? map estimates? pose and map estimates?
Active Localization [Fox et al., 1998] Active Mapping [Koenig et al., 2001] Active SLAM [Stachniss et al., 2005]	

Robotics: what is exploration?

Active Mapping (AM)

How to reduce map uncertainty ...

... when localization is reliable?

Related Work (View Candidates)

Distance	[Yamauchi, 1998, Bautin et al., 2012, Faigl et al., 2012]
Information	[González-Banos and Latombe, 2002, Burgard et al., 2005]
	[Amigoni and Gallo, 2005, Basilico and Amigoni, 2009]
	[Moorehead et al., 2001, Stachniss and Burgard, 2003]

Drawbacks

- ▶ Good performance in general but highly suboptimal in some environments
- ▶ Greedy evaluation from current map and position
- ▶ Few theoretical guarantees [Koenig et al., 2001, Tovey and Koenig, 2003]

Robotics: what about coverage?

Active Coverage (AC) [Galceran and Carreras, 2013]

How to fully observe the environment when localization is reliable and ...
... measurements are known?

Continuum of prior information



However, currently...

- ▶ Active Coverage techniques fail with incomplete information
- ▶ Active Mapping techniques are blind to prior information
- ▶ Active Coverage and Mapping objectives differ

→ We address both coverage and exploration in the SDM framework

Graph Theory: what is coverage?

Minimum Coverage Path (MCP) [Wernli, 2012, Garey and Johnson, 2002]

- ▶ Instance: Graph $G(V, E)$, a function $w_E: E \rightarrow \mathbb{Z}^+$, a source $s \in V$.
- ▶ Question: Find an **optimal coverage path**:
A sequence $S = \langle s = v_1, v_2, \dots, v_k \rangle$ such that $\forall i \in [1, k-1]$, $(v_i, v_{i+1}) \in E$ and
 $\bigcup_{i=1}^k v_i = V$, with minimum cost $\sum_{i=1}^{k-1} w_E(v_i, v_{i+1})$.
- ▶ Complexity: NP-hard

SDM formalism [Russell and Norvig, 1995, Li et al., 2012]

- ▶ Complete prior, deterministic control and sensing
- ▶ Completely Observable (State) Search Problem (COSP)

Graph Theory: what about exploration?

Canadian Traveler Problem (CTP) [Papadimitriou and Yannakakis, 1991, Fried et al., 2013]

- ▶ Instance: Graph $G(V, E)$, a function $w_E: E \rightarrow \mathbb{Z}^+$, a source $s \in V$, a target $t \in V$, a blocking probability $p: E \rightarrow [0, 1]$. Update blocking probability of neighboring edges upon node visitation.
- ▶ Question: Find an **optimal navigation policy**:
A mapping $\pi: V \times |E - 1|-\text{simplex} \rightarrow E$ that minimizes the expected traversal cost $C(\pi)$ from s to t .
- ▶ Complexity: PSPACE-hard

SDM formalism [Blei and Kaelbling, 1999, Nikolova and Karger, 2008]

- ▶ Incomplete prior, deterministic control and sensing
- ▶ Completely Observable (Belief state) Markov Decision Process (COMDP_{b^{az}})

→ Exploration is a Covering CTP! [Liao and Huang, 2014]

Graph Theory: illustration

	Truth	Known	Used	MCP_{\leftarrow}	MCP_{\rightarrow}	Decision
Coverage				$\Rightarrow \times 3$	$\Rightarrow, \leftarrow \times 3$	$\arg \min_{a \in \{\leftarrow, \rightarrow\}} \text{Cost}(MCP_a(p, m))$ where $\text{Cost}(MCP_{\leftarrow}(p, m)) = 4$ $\text{Cost}(MCP_{\rightarrow}(p, m)) = 5$ hence go $\leftarrow, \rightarrow \times 3$
Exploration				$\Rightarrow \times 3$	$\Rightarrow, \leftarrow \times 3$	
				$\Rightarrow \times 3$	$\Rightarrow, \leftarrow \times 2$	$\arg \min_{a \in \{\leftarrow, \rightarrow\}} \mathbb{E}_m[\text{Cost}(MCP_a(p, m))]$ where $\mathbb{E}_m[\text{Cost}(MCP_{\leftarrow}(p, m))] = 1 + 16/8$ $\mathbb{E}_m[\text{Cost}(MCP_{\rightarrow}(p, m))] = 1 + 18/8$ hence go \leftarrow

Summary

	Coverage	Exploration
Robotics	AC	AM
Graph Theory	MCP	CCTP
Sequential Decision-Making	COSP	COCP _{b^{az}}

- ▶ Study 1: Modeling coverage with knowledge rewards
- ▶ Study 2: Optimal coverage with search
- ▶ Study 3: Modeling exploration with belief rewards
- ▶ Study 4: Greedy exploration with interactions

Study 1: Models of coverage

Objective: model coverage with SDM formalisms

- ① Specify coverage domains
- ② Specify coverage models

	Truth	Known	Used	$MCP \Leftarrow$	$MCP \Rightarrow$	Decision
Coverage				$\Rightarrow \times 3$	$\Rightarrow, \Leftarrow \times 3$	$\arg \min_{a \in \{\Leftarrow, \Rightarrow\}} \text{Cost}(MCP_a(p, m))$ where $\text{Cost}(MCP \Leftarrow (p, m)) = 4$ $\text{Cost}(MCP \Rightarrow (p, m)) = 5$ hence go $\Leftarrow, \Rightarrow \times 3$

Assumptions

- ▶ Mono robot
- ▶ Static environments
- ▶ Stationary transitions
- ▶ Perfect sensing
- ▶ Complete prior

Background on SDM formalisms

Formalisms

- ▶ State
 - ▶ Actions
 - ▶ State-transition function
 - ▶ State-observation function
 - ▶ Rewards

State-transition function


- ▶ Deterministic: next state is known
 - ▶ Nondeterministic: unknown probability
 - ▶ Stochastic: known probability

Coverage with deterministic control and sensing

Deterministic control



World state, robot field of view raycasting, deterministic sensing



Formalism (Completely Observable (State) Search Problem (COSP))

$$\text{COSP} = \langle s_0, \text{Action}, \underline{\text{Transition}}, \text{Goal}, \text{Cost} \rangle$$


- ▶ s_0 , initial state
- ▶ Action(s), actions available from state s
- ▶ Transition(s, a), state attained from s by action a
- ▶ Goal(s), goal predicate for s
- ▶ Cost(s, a, s'), cost incurred from s to s' by a

Coverage with deterministic control and sensing

Deterministic control



World state, robot field of view raycasting, deterministic sensing



COSP-based coverage model: $\text{cov}_{det}(\text{map}: \langle 2, 2, m_0 \rangle, \text{robot}: \langle p_0, 0 \rangle)$

$$p_0 \in P = [1..rows] \times [1..cols]$$


- $s_0 \leftarrow \langle p_0, m_0^*, c_0 \rangle$ s.t. $m_0: P \rightarrow \{f \text{ (free)}, \neg f\}$
 $c_0: P \rightarrow \{c \text{ (covered)}, \neg c\}$

- $Action(s) \subseteq \{\uparrow, \Rightarrow, \downarrow, \Leftarrow\}$

- $Transition(s, a) \leftarrow \left\{ s' = \langle p', m', c' \rangle \mid \begin{array}{l} m' = nextOcc(m) \\ p' = nextPose(p, m', a) \\ c' = nextCov(p', m', c) \end{array} \right\}$


- $Goal(s) \leftarrow (s_c = c_G)$

- $Cost(s, a, s') \leftarrow \begin{cases} 0 & \text{if } Goal(s) \\ 1 - |s_c^{-1}(c) \setminus s_c^{-1}(c')| & \text{otherwise} \end{cases}$



Coverage with nondeterministic control and deterministic sensing

Nondeterministic control




Formalism (Completely Observable Contingency Problem (COPC))

$$\text{COPC} = \langle s_0, \text{Action}, \underline{\text{Transition}'}, \text{Goal}, \text{Cost} \rangle$$

- ▶ *Transition'(s,a)*, states attainable from s by action a
- ▶ *Other items as the COPC*


Coverage with nondeterministic control and deterministic sensing

Nondeterministic control




COCP-based coverage model: $cov_{nondet}(map:\langle 2,2,m_0 \rangle, robot:\langle p_0,0 \rangle)$

- ▶ $Transition'(s,a) \leftarrow \left\{ s' = \langle p', m', c' \rangle \mid \begin{array}{l} p' \in \text{nextPoses}(p, m', a) \\ m' = \text{nextOcc}(m) \\ c' = \text{nextCov}(p', m', c) \end{array} \right\}$
- ▶ Other items as cov_{det}



Coverage with stochastic control and deterministic sensing

Stochastic control




Formalism (Completely Observable Markov Decision Process (COMDP) [Puterman, 2014])

$$\text{COMDP} = \langle s_0, S, A, T, R \rangle$$

- ▶ s_0 , initial state
- ▶ S , finite set of states
- ▶ A , finite set of actions
- ▶ $T(s, a, s')$, probability of reaching s' from s by a
- ▶ $R(s, a, s')$, reward obtained from s to s' by a


Coverage with stochastic control and deterministic sensing

Stochastic control



COMDP-based coverage model: cov_{sto}

- ▶ $s_0 \leftarrow \langle p_0, m_0^*, c_0 \rangle$
- ▶ $S \leftarrow P \times M^* \times C$
- ▶ $A \leftarrow \{\uparrow, \Rightarrow, \Downarrow, \Leftarrow\}$
- ▶ $T(s, a, s') \leftarrow \Pr(s' = \langle p', m', c' \rangle | s = \langle p, m, c \rangle, a)$
 $= \Pr(m' | m) \times \Pr(p' | p, m', a) \times \Pr(c' | p', m', c)$
 $= \mathbb{1}_{\{m\}}(m') \times \textcolor{red}{\Pr(p' | p, m', a)} \times \mathbb{1}_{\{\text{nextCov}(p', m', c)\}}(c')$
- ▶ $R(s, a, s') \leftarrow -Cost(s, a, s')$



Conclusion

Models of coverage

We specified three knowledge state models with deterministic sensing and

- ① Deterministic control cov_{det}
- ② Nondeterministic control cov_{nondet}
- ③ Stochastic control cov_{sto}

Optimal coverage


- ① Complete prior
- ② Nondeterministic control
- ③ Deterministic sensing
- ④ Penalize time and reward information

→ Next, we address optimal coverage

Study 2: Planning for Deterministic Coverage

Objective: solving deterministic coverage with SDM planners

- ① Identify standard planning techniques
- ② Benchmark generic solvers and specific heuristics



Paradigms of deterministic search

Off-line [Dijkstra, 1959]

- ① Full plan
- ② Full execution

On-line [Knuth, 1977]

- ① Repair plan
- ② Full execution

Real-time [Korf, 1990]

- ① Partial plan
- ② Partial execution

Background on heuristic search

Core principles

- ▶ Generic solvers: optimal cost-so-far computation [Bellman, 1957]

$$g(s) \approx g^*(s) = \begin{cases} 0 & \text{if } \text{source}(n) \\ \min_{s' \in \text{Pred}(s)} g^*(s') + c(s', s) & \text{otherwise} \end{cases}$$

- ▶ Specific heuristic: optimal cost-to-go estimate

$$h(s) \approx h^*(s)$$

Algorithms and heuristics

Generic Solvers

Algorithm	Selection for expansion	f-Evaluation	g-Optimality
Best-First Search Framework			
Iterative Best-First Search (I-BFS) [Dechter and Pearl, 1985]	f-promising Open node	agg(g, h)	Yes*
Iterative A* (I-BFS(A*)) [Hart et al., 1968, Hart et al., 1972]	—	g + h	Yes*
Iterative Weighted A* (I-BFS(wA*)) [Pohl, 1970]	—	g + wh	Yes*
Frontier Best-First Search (F-BFS) [Korf et al., 2005]	—	agg(g, h)	Yes*
Recursive Best-First Search (R-BFS) [Korf, 1993]	F-promising subtree root	agg(g, h)	Yes*
Depth-First Search Framework			
Depth-First Search (I-DFS, R-DFS)	trajectory end-node successor	agg(g, h)	No
Iterative-Deepening Depth-First Search (ID-DFS) [Korf, 1985]	s.t. $f(\text{succ}) \leq f\text{-lowerbound}$	agg(g, h)	Yes*
Branch-and-Bound Depth-First Search (BnB-DFS) [Zhang and Korf, 1995]	s.t. $f(\text{succ}) < f\text{-upperbound}$	agg(g, h)	Yes*

Specific heuristics

Coverage: What is not yet covered?

$$\text{MaxCells}(s) = \left\lceil \frac{|C_A^{\neg c}(s)|}{\max_{p,p' \in P} |F(p', s_m) \setminus F(p, s_m)|} \right\rceil$$

$$C_A^{\neg c}(s) = \{p \in C_A^{\max} \mid s_c(p) = \neg c\}$$

Navigation: How far is the information?

$$\text{MinDist}(s) = \min_i \min_{v \in V_i(s)} d^*(s_p, v)$$

$$\text{MaxDist}(s) = \max_i \min_{v \in V_i(s)} d^*(s_p, v)$$

$$V_i(s) = \{p \in P_A^{\max} \mid F(s_p, s_m) \cap C_A^{\neg c, i}(s) \neq \emptyset\}$$


Experimental protocol

Questions


- ① What can we solve optimally with few resources?
 - ② Otherwise, can we still obtain some guarantees?

Parameters		
Instances	empty map not empty map	291 instances 300 instances
Solvers	BFS DFS	I-BFS, R-BFS ID-DFS, BnB-DFS
Heuristics	coverage navigation	MaxCells MinDist, MaxDist
Metrics	solution cost memory consumption time consumption instance difficulty	path length node storage node generation <i>branching</i> ^{depth}
Resources	memory out time out	1.10^6 nodes 2.10^7 nodes
Details	action ordering second tie-breaker	$\{1^{\text{st}} \uparrow, \rightarrow, \downarrow, \leftarrow\}$ $\text{lexmin } f(\cdot), h(\cdot)$

#obstacles	#rows	#columns	#starts	#instances
0	1	[1-25]	[1-13]	169
0	2	[2-13]	[1-7]	48
0	3	[3-9]	[3-10]	45
0	4	[4-7]	[3-8]	23
0	5	5	6	6
5	5	5	6	300




Results: optimal solvers $f(n)=g(n)+h(n)$



Coverage heuristics orders of magnitude faster

Results: anytime solver BnB-DFS(wA*)



Conclusion

On optimal coverage $\text{Cost}(MCP)$

- ➊ Optimal: coverage \gg navigation heuristics
- ➋ Suboptimal: guarantees ϵ -optimal solutions
- ➌ Anytime: converges to ϵ -optimal coverage

On optimal exploration $\mathbb{E}_m[\text{Cost}(MCP)]$

- ➊ Coverage pattern database
- ➋ Optimally solve exploration

→ Next, we model exploration with an incomplete prior

Study 3: Models of exploration

Objective: Model exploration with SDM formalisms

- ① Specify exploration domains
- ② Specify exploration models

	Truth	Known	Used	MCP_{\Leftarrow}	MCP_{\Rightarrow}	Decision
Exploration				$\Rightarrow \times 3$	$\Leftarrow, \Leftarrow \times 3$	
				$\Rightarrow \times 3$	$\Rightarrow, \Leftarrow \times 2$	$\arg\min_{a \in \{\Leftarrow, \Rightarrow\}} \mathbb{E}_m[\text{Cost}(MCP_a(p, m))]$
				$\Rightarrow \times 2$	$\Leftarrow \times 1$	$\mathbb{E}_m[\text{Cost}(MCP_{\Leftarrow}(p, m))] = 1 + 16/8$
				$\Rightarrow \times 2$	$\Leftarrow \times 1$	$\mathbb{E}_m[\text{Cost}(MCP_{\Rightarrow}(p, m))] = 1 + 18/8$
				$\Rightarrow \times 2$	$\Rightarrow, \Leftarrow \times 3$	
				$\Rightarrow \times 2$	$\Rightarrow, \Leftarrow \times 2$	
				$\Rightarrow \times 1$	$\Leftarrow \times 1$	
				$\Rightarrow \times 1$	$\Leftarrow \times 1$	
						hence go \Leftarrow

Assumptions

- ▶ Imperfect sensing
- ▶ Incomplete prior

Background on SDM formalisms

Formalisms

- ▶ State
- ▶ Actions
- ▶ State-transition function
- ▶ State-observation function
- ▶ Rewards

State-observation function



- ▶ Deterministic: single measurement with certainty
- ▶ Nondeterministic: unknown probability
- ▶ Stochastic: known probability

Domain of exploration with deterministic control and sensing

Deterministic control



World state, robot field of view raycasting, expected view



Formalism (Partially Observable (State) Search Problem (POSP))

$$\text{POSP} = \langle s_0, \text{Action}, \text{Transition}, \text{Goal}, \text{Cost}, \underline{\text{Observation}}^{(')} \rangle$$


- ▶ $\langle s_0, \text{Action}, \text{Transition}, \text{Goal}, \text{Cost} \rangle$ is a COSP
- ▶ $\text{Observation}^{(')}(s, a, s')$, potential measurements from s' after applying a from s

Domain of exploration with deterministic control and sensing

Deterministic control



World state, robot field of view raycasting, expected view




POSP-based domain of exploration: \exp_{det}


- ▶ $s_0 \leftarrow \langle p_0, m_0 \rangle$
- ▶ $Action(s) \subseteq \{noop, \uparrow, \Rightarrow, \downarrow, \Leftarrow\}$
- ▶ $Transition(s, a) \leftarrow \left\{ s' = \langle p', m' \rangle \mid \begin{array}{l} m' = nextOcc(m) \\ p' = \text{nextPose}(p, m', a) \end{array} \right\}$
- ▶ $Goal(\cdot) \leftarrow False$
- ▶ $Cost(a) \leftarrow 1 + \mathbb{1}_{\neq noop}(a)$
- ▶ $Observation(s) \leftarrow \left\{ z_s = \langle z_p, z_v \rangle \mid \begin{array}{l} z_p = s_p \\ z_v = \text{expectedView}(s_p, s_m) \end{array} \right\}$

Domain of exploration with nondeterministic control and sensing

Nondeterministic control



Nondeterministic sensing ([0..1]-error)




Formalism (Partially Observable Contingency Problem (POCP))

$$\text{POCP} = \langle s_0, \text{Action}, \text{Transition}', \text{Goal}, \text{Cost}, \underline{\text{Observation}'} \rangle$$


- ▶ $\langle s_0, \text{Action}, \text{Transition}', \text{Goal}, \text{Cost} \rangle$ is a COCP
- ▶ $\text{Observation}'(s, a, s')$, potential measurements from s' after applying a from s

Domain of exploration with nondeterministic control and sensing

Nondeterministic control



Nondeterministic sensing ([0..1]-error)




POCP-based domain of exploration: \exp_{nondet}

- ▶ $Transition'(s, a) \leftarrow \left\{ s' = \langle p', m' \rangle \mid \begin{array}{l} m' = nextOcc(m) \\ p' \in nextPoses(p, m', a) \end{array} \right\}$
- ▶ $Observation'(s) \leftarrow \left\{ z_s = \langle z_p, z_v \rangle \mid \begin{array}{l} z_p = s_p \\ z_v \in potentialViews(s_p, s_m, error_{max}) \end{array} \right\}$
- ▶ Other items as \exp_{det}

Domain of exploration with stochastic control and sensing

Stochastic sensing




Formalism (Partially Observable Markov Decision Process (POMDP) [Littman, 1996])

$$\text{POMDP} = \langle s_0, S, A, T, R, \underline{Z}, \underline{Q} \rangle$$

- ▶ $\langle s_0, S, A, T, R \rangle$ is a COMDP
- ▶ Z , finite set of measurements
- ▶ $O(a, s', z)$, probability of measuring z from state s' after applying a


Domain of exploration with stochastic control and sensing

Stochastic sensing



POMDP-based domain of exploration: \exp_{sto}

- ▶ $s_0 : \langle p_0, m_0 \rangle$
- ▶ $S \leftarrow P \times M$
- ▶ $A \subseteq \{\text{noop}, \uparrow, \Rightarrow, \downarrow, \Leftarrow\}$
- ▶ $T(s, a, s') \leftarrow \mathbb{1}_{\{m\}}(m') \times \Pr(p' | p, m', a)$
- ▶ $R(a) \leftarrow -Cost(a)$
- ▶ $Z \leftarrow P \times V$
- ▶ $O(s, z) \leftarrow \Pr(z = \langle p_z, v_z \rangle | s = \langle p, m \rangle)$
 $= \Pr(p_z | p) \times \Pr(v_z | v_{exp})$
 $= \mathbb{1}_{\{p\}}(p_z) \times \frac{score(v, v_{exp})}{\sum_{v_z \in V(s)} score(v_z, v_{exp})}$



From external domains to internal models

Domains of exploration

We specified three external state domains to simulate exploration

- ① Deterministic control and sensing exp_{det}
- ② Nondeterministic control and sensing exp_{nondet}
- ③ Stochastic control and sensing exp_{sto}

Models of exploration with belief-dependent rewards

However, the agent must rely on internal belief models to handle

- ① Nondeterministic sensing
- ② Incomplete prior

Model of exploration with nondeterministic control and sensing

Formalism (Completely Observable (Belief) Contingency Problem ($COCP_{baz}$))

$$COCP_{baz} = \langle b_0, ApplicableAction_b, TransitionModel'_b, GoalTest_b, StepCost'_b \rangle$$

- ▶ b_0 , initial belief
- ▶ $Action_b(b)$, actions available from belief b
- ▶ $Transition'_b(b,a)$, beliefs attainable from b by a
- ▶ $Goal_b(b)$, goal predicate for b
- ▶ $Cost'_b(b,a,z)$, cost incurred from b after a and z

Model of exploration with nondeterministic control and sensing

COCP_{b^{az}}-based model of exploration: $\exp_{(non)det}$


- ▶ $b_0 \leftarrow P_0 \times M_0$
 $P_0 = \{p_0 = \langle 1, 1 \rangle\}$

$$M_0 = \left\{ m_i \mid m_i(p) \in \begin{cases} \{f\} & \text{if } p \in \{p_0\} \\ \{f, \neg f\} & \text{otherwise} \end{cases} \right\}$$

- ▶ $Action_b(b) \leftarrow \bigcup_{s \in b} Action(s)$
- ▶ $Goal_b \leftarrow |b| = 1$
- ▶ $Transition'_b(b, a) \leftarrow \{b^{az} \mid \forall z \in Z^a\}$
- ▶ $Z^a = \bigcup_{s' \in b^a} Observation^{(')}(s')$
- ▶ $b^a = \bigcup_{s \in b} Transition^{(')}(s, a)$
- ▶ $b^{az} = \{s' \in b^a \mid z \in Observation^{(')}(s')\}$
- ▶ $Cost'_b(b, a, z) \leftarrow Cost(a) - |b \setminus b^{az}| (\rho \text{POSP} / \rho \text{POCP})$

Model of exploration with nondeterministic control and sensing

COCP_{baz}-based belief space of exploration: \exp_{det}



Model of exploration with stochastic control and sensing

Formalism (Completely Observable (Belief state) Markov Decision Process ($COMDP_{b^{az}}$))

$$COMDP_{b^{az}} = \langle b_0, B, A_b, T'_b, R'_b, \rangle$$

- ▶ b_0 , initial belief state
- ▶ B , continuous belief-state space
- ▶ A_b , finite set of actions
- ▶ $T'_b(b, a, b')$, probability of reaching b' from b by a
- ▶ $R'_b(b, a, z)$, reward from b after a and z

Model of exploration with stochastic control and sensing


COMDP_{b^{az}}-based model of exploration: exp_{sto}

- ▶ $b_0(s) = P(S_0 = s) = \mathbf{U}_{S_0}(\{s\})$
- ▶ $B = |S - 1| - \text{simplex}$
- ▶ $A_b(b) = \bigcup_{s \in S, b(s) > 0} A(s)$
- ▶ $T_b(b, a, b') = Pr(b' | b, a) = \sum_{z \in Z} \mathbb{1}_{\{b^{az}\}}(b') Pr(z | b, a)$
 $b^a(s') = Pr(s' | b, a) = \sum_{s \in S} b(s) T(s, a, s')$

$$b^{az}(s') = Pr(s' | b, a, z) = \eta O(s', z) b^a(s')$$

$$Pr(z | b, a) = \frac{1}{\eta} = \sum_{s' \in S} O(s', z) b^a(s')$$

- ▶ $R'_b(b, a, z) = R(a) + D_{KL}(b^{az} || b)$ (ρ POMDP [Araya-López et al., 2010])



Conclusion

Models of exploration

We reported two models with belief-dependent rewards

- ① Unquantified beliefs $\exp_{(non)det}$
- ② Quantified beliefs \exp_{sto}

Optimal exploration

- ① Incomplete prior
- ② Nondeterministic control and sensing
- ③ Action penalty and information reward

Unfortunately,...

- ① Pioneering algorithms do not scale [Araya-López et al., 2010]
- ② Current implementation solve small instances (6 cells)
- ③ State-of-the-art algorithms do not reward beliefs [Silver and Veness, 2010]

→ Next, we address greedy exploration

Study 4: Multi-Robot Exploration in Dynamic Environments

Objective: exploring crowded indoor environments

- ➊ Robots jointly build a map
- ➋ Indoor crowded environments
- ➌ Exploit pedestrian flows



Assumptions

- ▶ Multi robot (decentralized)
- ▶ Dynamic environment
- ▶ Stationary transitions
- ▶ Perfect sensing
- ▶ Incomplete prior

Background on greedy exploration

Multi-robot exploration


- ▶ fixed:
Hardwired [Baronov and Baillieul, 2007, Morlok and Gini, 2007]
- ▶ reactive:
Ants [Koenig and Liu, 2001, Svennebring and Koenig, 2004, Ferranti et al., 2007, Glad et al., 2010, Andries and Charpillet, 2013]
- ▶ goal-based:
MinDist [Yamauchi, 1997], MinPos [Bautin et al., 2012], MTSP [Faigl et al., 2012]
- ▶ utility-based:
 $V - \beta C_w$ [Burgard et al., 2005], $Ae^{-\lambda C}$ [González-Banos and Latombe, 2002], Pareto [Amigoni and Gallo, 2005], MCDM [Basilico and Amigoni, 2009], $\alpha \mathbb{E}[I] - C$ [Stachniss and Burgard, 2003]
- ▶ others:
Auction [Zlot et al., 2002], Motivation [Macedo and Cardoso, 2004]

Multi-robot crowded exploration domain

World state

$m: P \rightarrow \{f, \neg f_s\}$, map


- $s: \langle m, \mathbf{R}, \mathbf{H} \rangle$ s.t. $\mathbf{R} = \{R_1, \dots, R_n\}$, robots
 $\mathbf{H} = \{H_1, \dots, H_m\}$, pedestrians



Robot measurement


$v_Z: FOV_i \rightarrow \{f, \neg f_s, \neg f_m, \neg k\}$

- $Z_i: \langle v_z, R_Z, H_Z \rangle$ s.t. $R_Z \subseteq \mathbf{R}$, visible robots
 $H_Z \subseteq \mathbf{H}$, visible pedestrians



Robot communication

- $\Theta_i^{0:t} = \bigcup_{j=1}^{|R|} Z_j^{0:t_j}$ s.t. $Z_i^{0:t}$, robot history



Multi-robot crowded exploration model

Robot knowledge map

- $k_i: P \rightarrow \{f, \neg f_s, \neg f_m, \neg k\}$

R_1			
H_1			H_2
			R_2

Targets T

- Frontiers F and Interactions H

R_1		F		
	F	F		F
				R_2


R_1				
	H			H
				R_2

On-line Task-Allocation Framework


$$\begin{array}{|c|c|} \hline c_{R_i T_j} & \mathbf{T} \\ \hline \mathbf{R} & \mathbf{C}_{\mathbf{RT}} \\ \hline \end{array} \xrightarrow{\text{opt.}} \begin{array}{|c|c|} \hline a_{R_i T_j} & \mathbf{T} \\ \hline \mathbf{R} & \mathbf{A}_{\mathbf{RT}} \\ \hline \end{array}$$

Robot-Target assignment costs C_{RT}

- $C_{RT} = \alpha \cdot D_{RT} + (1 - \alpha) \cdot P_{RT}, \alpha \in [0, 1]$



$$d_{R_i X_j} = \text{shortest_path_length}_{R_i, X_j}$$




$$p_{R_i X_j} = \text{idle_time}_{R_i X_j} + \text{reorientation}_{R_i X_j}$$

Experimental protocol

Parameters

Maps	empty unstructured structured	10×10 cells ($25m^2$) 12×12 cells ($36m^2$) 22×11 cells ($60.5m^2$)
Strategies	greedy	local, group
Pedestrian	speed controller starting	$0.56m/s$ V-REP Walking Bill uniform random
Robot	speed viewing range network range starting	$0.5m/s$ $2m$ $100m$ fixed positions
Modulators	α σ	$0:0.25:1$ $0:0.25:1$
Metrics	coverage distance time	$\frac{\text{covered_cells}}{\text{all_cells}}$ m s
Resources	time out	300s

Maps



Greedy strategies

Group: And the others?


Local: Only me

- ① Initialize $\mathbf{T} \leftarrow \{\mathbf{H} \cup \mathbf{F}\}$
- ② Find $T^* \leftarrow \arg\min_{T_j \in \mathbf{T}} CR_i T_j$


- ① Initialize $\mathbf{T} \leftarrow \{\mathbf{H} \cup \mathbf{F}\}; R_V \leftarrow R_Z$
- ② Find $(R^*, T^*) \leftarrow \arg\min_{R_j \in R_V, T \in \mathbf{T}} CR_j T_j$
- ③ If $(R^* = R_i)$ then return T^* ; else $R_Z \leftarrow R_Z \setminus R^*; \mathbf{T} \leftarrow \mathbf{T} \setminus T^*$
- ④ Go to step 2 unless \mathbf{T} is empty

Results: density (30%), robot average (20 samples)

Local Greedy



Group Greedy



Best parameters (0,0);(0.75,0);(0.5,0)

Best parameters (0.25,0);(0.5,0);(0.25,0)

frontiers are never penalized but pedestrians are

Conclusion

Hybrid Exploration

- ① Enabled the frontier paradigm in crowded environment
- ② Proposed an interactive paradigm to exploit affordances

Early Results

- ① Targets: global frontiers and local interactions
- ② Costs: weighted distance and time-orientation penalty
- ③ Static tuning: best results when interactions are penalized

→ Finally, we conclude

Summary

Studies 1 and 3: Modeling coverage and exploration with SDM formalisms

- ▶ Lookahead models with knowledge/belief rewards
- ▶ From complete to incomplete prior
- ▶ From deterministic to stochastic control and observation
- ▶ Allow **optimal centralized coverage and exploration**

Study 2: Optimally solving coverage with SDM planners

- ▶ Benchmark of generic solvers and expert heuristics
- ▶ Orders of magnitude domination of coverage heuristics
- ▶ Anytime improvement of exploration trajectories

Study 4: Greedy exploration with interactions

- ▶ Crowded environments
- ▶ Hybrid exploration paradigm
- ▶ Frontiers preferred over interactions

Perspectives

Studies 1 and 3: Modeling coverage and exploration with SDM formalisms

- ▶ Centralized multi-robot: ρ POMDP [Araya-López et al., 2010]
- ▶ Decentralized multi-robot: Dec- ρ POMDP [Renoux, 2015]

Study 2: Optimally solving coverage with SDM planners

- ▶ Evaluate other heuristics: MTSP [Faigl et al., 2012]
- ▶ Consider other paradigms: On-Line [Knuth, 1977], Real-Time search [Korf, 1990]

Study 4: Greedy exploration with interactions

- ▶ Realistic pedestrian behavior: Social Force Model [Helbing and Molnar, 1995]
- ▶ On-line parameter tuning: Embodied Evolution [Watson et al., 2002]

The End

Thank you for your attention. Questions?

List of contributions

- ① Models of coverage and exploration
- ② Planning for optimal coverage
- ③ Greedy exploration in crowds

References |

-  Amigoni, F. and Gallo, A. (2005).
A Multi-Objective Exploration Strategy for Mobile Robots.
In *Proceedings of the International Conference on Robotics and Automation*.
-  Andries, M. and Charpillet, F. (2013).
Multi-Robot Exploration of Unknown Environments with Identification of Exploration Completion and Post-Exploration Rendez-Vous using Ant Algorithms.
In *Proceedings of the International Conference on Intelligent Robots and Systems*.
-  Araya-López, M., Buffet, O., Thomas, V., and Charpillet, F. (2010).
A POMDP Extension with Belief-Dependent Rewards.
In *Advances in Neural Information Processing Systems*.
-  Bajcsy, R. (1988).
Active Perception.
Proceedings of the IEEE, 76(8).

References II

-  Baronov, D. and Baillieul, J. (2007).
Reactive Exploration through Following Isolines in a Potential Field.
In *Proceedings of the American Control Conference*.
-  Basilico, N. and Amigoni, F. (2009).
Exploration Strategies based on Multi-Criteria Decision Making for an Autonomous Mobile Robot.
In *Proceedings of the European Conference on Mobile Robots*, pages 259–264.
-  Bautin, A., Simonin, O., and Charpillet, F. (2012).
MinPos : A Novel Frontier Allocation Algorithm for Multi-robot Exploration.
In *Proceedings of the International Conference on Intelligent Robotics and Applications*.
-  Bellman, R. (1957).
Dynamic Programming.
Princeton University Press.

References III

-  Blei, D. M. and Kaelbling, L. P. (1999).
Shortest Paths in a Dynamic Uncertain Domain.
In *IJCAI Workshop on Adaptive Spatial Representations of Dynamic Environments*, volume 4, page 2.
-  Burgard, W., Moors, M., Stachniss, C., and Schneider, F. (2005).
Coordinated Multi-Robot Exploration.
Transactions on Robotics, 21(3).
-  Dechter, R. and Pearl, J. (1985).
Generalized Best-First Search Strategies and the Optimality of A*.
Journal of the ACM, 32(3).
-  Dijkstra, E. (1959).
A Note on Two Problems in Connexion with Graphs.
Numerische mathematik, 1(1).
-  Durrant-Whyte, H. and Bailey, T. (2006).
Simultaneous Localization and Mapping: Part I.
Robotics & Automation Magazine, IEEE, 13(2).

References IV

-  Elfes, A. (1989).
Using Occupancy Grids for Mobile Robot Perception and Navigation.
Computer, 22(6).
-  Faigl, J., Kulich, M., and Přeučil, L. (2012).
Goal Assignment using Distance Cost in Multi-Robot Exploration.
In *Proceedings of the International Conference on Intelligent Robots and Systems*.
-  Ferranti, E., Trigoni, N., and Levene, M. (2007).
Brick&Mortar: an On-Line Multi-Agent Exploration Algorithm.
In *Proceedings of the International Conference on Robotics and Automation*.
-  Fox, D., Burgard, W., and Thrun, S. (1998).
Active Markov Localization for Mobile Robots.
Robotics and Autonomous Systems, 25(3-4).

References V

-  Fried, D., Shimony, S. E., Benbassat, A., and Wenner, C. (2013).
Complexity of Canadian Traveler Problem Variants.
Theoretical Computer Science, 487:1–16.
-  Galceran, E. and Carreras, M. (2013).
A Survey on Coverage Path Planning for Robotics.
Robotics and Autonomous Systems, 61(12).
-  Garey, M. and Johnson, D. (2002).
Computers and Intractability, volume 29.
-  Glad, A., Simonin, O., Buffet, O., and Charpillet, F. (2010).
Influence of Different Execution Models on Patrolling Ant Behaviors: from Agents to Robots.
In *International Conference on Autonomous Agents and Multiagent Systems*.
-  González-Banos, H. H. and Latombe, J.-C. (2002).
Navigation Strategies for Exploring Indoor Environments.
The International Journal of Robotics Research, 21(10-11).

References VI

-  Hart, P., Nilsson, N., and Raphael, B. (1968).
A Formal Basis for the Heuristic Determination of Minimum Cost Paths.
Transactions on Systems Science and Cybernetics, 4(2).
-  Hart, P., Nilsson, N., and Raphael, B. (1972).
Correction to a Formal Basis for the Heuristic Determination of Minimum Cost Paths.
ACM SIGART Bulletin, (37).
-  Helbing, D. and Molnar, P. (1995).
Social Force Model for Pedestrian Dynamics.
Physical review E, 51(5).
-  Knuth, D. (1977).
A Generalization of Dijkstra's Algorithm.
Information Processing Letters, 6(1).
-  Koenig, S. and Liu, Y. (2001).
Terrain Coverage with Ant Robots: A Simulation Study.
In *Proceedings of the international conference on Autonomous agents*.

References VII

-  Koenig, S., Tovey, C., and Halliburton, W. (2001).
Greedy Mapping of Terrain.
In Proceedings of the International Conference on Robotics and Automation, volume 4.
-  Korf, R. (1985).
Depth-First Iterative-Deepening: An Optimal Admissible Tree Search.
Artificial intelligence, 27(1).
-  Korf, R. (1990).
Real-Time Heuristic Search.
Artificial intelligence, 42(2).
-  Korf, R. (1993).
Linear-Space Best-First Search.
Artificial Intelligence, 62(1).
-  Korf, R., Zhang, W., Thayer, I., and Hohwald, H. (2005).
Frontier Search.
Journal of the ACM, 52(5).

References VIII

-  LaValle, S. M. (2006).
Planning Algorithms.
Cambridge university press.
-  Li, A. Q., Amigoni, F., and Basilico, N. (2012).
Searching for Optimal Off-Line Exploration Paths in Grid Environments for
a Robot with Limited Visibility.
In *AAAI*.
-  Liao, C.-S. and Huang, Y. (2014).
The Covering Canadian Traveller Problem.
Theoretical Computer Science, 530(Supplement C):80 – 88.
-  Littman, M. (1996).
Algorithms for Sequential Decision Making.
PhD thesis, Brown University.

References IX

-  Macedo, L. and Cardoso, A. (2004).
Exploration of Unknown Environments with Motivational Agents.
In *Proceedings of the International Joint Conference on Autonomous Agents and Multiagent Systems*.
-  Mihaylova, L., Lefebvre, T., Bruyninckx, H., Gadeyne, K., and De Schutter, J. (2003).
A Comparison of Decision Making Criteria and Optimization Methods for Active Robotic Sensing.
In *Numerical Methods and Applications*.
-  Moorehead, S. J., Simmons, R., and Whittaker, W. L. (2001).
Autonomous Exploration using Multiple Sources of Information.
In *Proceedings of the International Conference on Robotics and Automation*, volume 3.
-  Morlok, R. and Gini, M. (2007).
Dispersing Robots in an Unknown Environment.
In *Distributed Autonomous Robotic Systems*.

References X

-  Nikolova, E. and Karger, D. (2008).
Route Planning under Uncertainty: The Canadian Traveller Problem.
In *AAAI*.
-  Papadimitriou, C. and Yannakakis, M. (1991).
Shortest Paths without a Map.
Theoretical Computer Science, 84(1).
-  Pohl, I. (1970).
Heuristic search viewed as path finding in a graph.
Artificial Intelligence, 1(3).
-  Puterman, M. (2014).
Markov Decision Processes: Discrete Stochastic Dynamic Programming.
-  Renoux, J. (2015).
Contribution to Multiagent Planning for Active Information Gathering.
PhD thesis.
-  Russell, S. and Norvig, P. (1995).
Artificial Intelligence: Modern Approach.

References XI

-  Silver, D. and Veness, J. (2010).
Monte-Carlo Planning in Large POMDPs.
In *Advances in neural information processing systems*.
-  Stachniss, C. and Burgard, W. (2003).
Exploring Unknown Environments with Mobile Robots using Coverage Maps.
In *Proceedings of the International Joint Conference on Artificial Intelligence*.
-  Stachniss, C., Grisetti, G., and Burgard, W. (2005).
Information Gain-based Exploration Using Rao-Blackwellized Particle Filters.
In *Robotics: Science and Systems*, volume 2.
-  Svennebring, J. and Koenig, S. (2004).
Building Terrain-Covering Ant Robots: A Feasibility Study.
Autonomous Robots, 16(3).


References XII

-  Thrun, S., Fox, D., Burgard, W., and Dellaert, F. (2001).
Robust Monte Carlo Localization for Mobile Robots.
Artificial Intelligence, 128(1-2).
-  Tovey, C. and Koenig, S. (2003).
Improved Analysis of Greedy Mapping.
In *Proceedings of the International Conference on Intelligent Robots and Systems*, volume 4.
-  Watson, R. A., Ficici, S. G., and Pollack, J. B. (2002).
Embodied Evolution: Distributing an Evolutionary Algorithm in a Population of Robots.
Robotics and Autonomous Systems, 39(1).
-  Wernli, D. (2012).
Grid Exploration.
Master's thesis, ETH Zürich, Department of Computer Science.


References XIII

-  Yamauchi, B. (1997).
A Frontier-Based Approach for Autonomous Exploration.
In *Proceedings of the International Symposium on Computational Intelligence in Robotics and Automation*.
-  Yamauchi, B. (1998).
Frontier-based Exploration using Multiple Robots.
In *Proceedings of the International Conference on Autonomous agents*.
-  Zhang, W. and Korf, R. (1995).
Performance of Linear-Space Search Algorithms.
Artificial Intelligence, 79(2).
-  Zlot, R., Stentz, A., Dias, M. B., and Thayer, S. (2002).
Multi-robot Exploration Controlled by a Market Economy.
In *Proceedings of the International Conference on Robotics and Automation*, volume 3.

Results: optimal solvers $f(n)=g(n)+h(n)$ 1/2




MinDist < MaxDist < MaxCells < MinDist + MaxCells





Orders of magnitude faster Coverage >> Navigation

Results: optimal solvers $f(n)=g(n)+h(n)$ 2/2





Results: suboptimal solvers $f(n)=g(n)+wh(n)$ 1/2



Monotonic I-BFS BnB-DFS, Nonmonotonic
ID-DFS, R-BFS

1-opt in theory → better in practice

Results: suboptimal solvers $f(n)=g(n)+wh(n)$ 2/2



Results: anytime solver BnB-DFS(wA*)

Opt convergence → 2-Opt convergence (10^2 faster)

